Thomas Hartvigsen
Cambridge, MA 02138
https://thartvigsen.github.io
tomh@mit.edu

Research Interests

Machine Learning, Data Mining, Time Series, Natural Language Processing, Healthcare, Fairness in AI systems

Employment

MIT 01/2022 - present
Postdoctoral Associate at CSAIL with Marzyeh Ghassemi

Worcester Polytechnic Institute 08/2016 - 12/2021
Research Fellow with Elke Rundensteiner and Xiangnan Kong

Microsoft 05/2021 - 08/2021
PhD Intern with Dipankar Ray and Hamid Palangi

UMass Medical School 08/2018 - 09/2019
Research Intern with Jomol Matthew

University of Arizona 05/2015 - 08/2015
NSF REU Intern with Shirley Papuga

Education

Worcester Polytechnic Institute, *Worcester, MA* 08/2016 - 12/2021
PhD, Data Science
MS, Data Science
Advised by Elke Rundensteiner and Xiangnan Kong

SUNY Geneseo, *Geneseo, NY*
BA, Applied Mathematics, minor in Biomathematics 08/2012 - 05/2016

Grants

NSF-III: Timely Classification for Actionable Predictions (Under Review)
PI: Elke Rundensteiner, Co-PI: Xiangnan Kong.
This grant proposal is written based on my research and I am responsible for 90% of the writing.

Selected Honors & Awards

🏆 **Best Poster**, International Conference on Health Informatics 2020
🏆 **Outstanding Graduate Research Award**, WPI 2019
🏆 **Best Poster**, Graduate Research Innovation and Exchange, WPI 2019
IMA Travel Award, University of Minnesota 2019
🏆 **People's Choice Poster Award**, Graduate Research Innovation and Exchange, WPI 2017
GAANN Fellowship (Annual Tuition + Stipend Award), U.S. Dept. of Education 2016-2021

Publications

I have published in KDD, AAAI, ACL, NeurIPS, CIKM, SDM, ECML, BigData, HEALTHINF, and BHI.

Refereed

21. ToxiGen: *Controlling Language Models to Generate Implied and Adversarial Toxicity*.
Thomas Hartvigsen, Saadia Gabriel, Hamid Palangi, Maarten Sap, Dipankar Ray, Ece Kamar.
ACL, 2022.

20. *Recovering the Propensity Score from Biased Positive Unlabeled Data*.
Walter Gerych, **Thomas Hartvigsen**, Luke Buquicchio, Emmanuel Agu, Elke Rundensteiner.
AAAI, 2022. **Oral Spotlight**.

19. *Positive Unlabeled Learning with a Sequential Selection Bias*.
Walter Gerych, **Thomas Hartvigsen**, Luke Buquicchio, Kavin Chandrasekaran, Abdulaziz Alajaji, Hamid Mansoor, Elke Rundensteiner, Emmanuel Agu.
SDM, 2022.

18. *Recurrent Bayesian Classifier Chains for Exact Multi-label Classification*.
Walter Gerych, **Thomas Hartvigsen**, Luke Buquicchio, Emmanuel Agu, Elke Rundensteiner.
NeurIPS, 2021.

17. *Energy-Efficient Models for High-Dimensional Spike Train Classification using Sparse Spiking Neural Networks*.
Hang Yin, John Boaz Lee, Xiangnan Kong, **Thomas Hartvigsen**, Sihong Xie.
KDD, 2021.

16. *Semi-Supervised Knowledge Amalgamation for Sequence Classification*.
Jidapa Thadajarassiri, **Thomas Hartvigsen**, Xiangnan Kong, Elke Rundensteiner.
AAAI, 2021.

15. *Learning Saliency Maps to Explain Deep Time Series Classifiers*.
Prathyush Parvatharaju, Ramesh Doddaiah, **Thomas Hartvigsen**, Elke Rundensteiner.
CIKM, 2021.

14. *Variational Open-Set Recognition*.
Luke Buquicchio, Walter Gerych, Kavin Chandrasekaran, Abdulaziz Alajaji, Hamid Mansoor, **Thomas Hartvigsen**, Elke Rundensteiner, Emmanuel Agu.
IEEE BigData, 2021.

13. *Human-like Explanation for Text Classification with Limited Attention Supervision*.
Dongyu Zhang, Cansu Sen, Jidapa Thadajarassiri, **Thomas Hartvigsen**, Xiangnan Kong, Elke Rundensteiner.
IEEE BigData, 2021.

12. *Recurrent Halting Chain for Early Multi-label Classification*.
Thomas Hartvigsen, Cansu Sen, Xiangnan Kong, Elke Rundensteiner.
KDD, 2020.

11. *Human Attention Maps for Text Classification: Do Humans and Neural Networks Focus on the Same Words?*
Cansu Sen, **Thomas Hartvigsen**, Biao Yin, Xiangnan Kong, Elke Rundensteiner.
ACL, 2020.

10. *Learning to Selectively Update State Neurons in Recurrent Networks*.
Thomas Hartvigsen, Cansu Sen, Xiangnan Kong, Elke Rundensteiner.
CIKM, 2020.

9. *Learning Similarity-Preserving Word Meta-Embedding*.
Jidapa Thadajarassiri, Cansu Sen, **Thomas Hartvigsen**, Xiangnan Kong, Elke Rundensteiner.
IEEE BigData, 2020.

8. *Clinical Performance Evaluation of a Machine Learning System for Predicting Hospital-Acquired Clostridium Difficile Infection*.
Erin Teeple, **Thomas Hartvigsen**, Cansu Sen, Kajal Claypool, Elke Rundensteiner.
HEALTHINF, 2020. 🏆 **Best Poster**.

7. *Adaptive-Halting Policy Network for Early Classification*.
Thomas Hartvigsen, Cansu Sen, Xiangnan Kong, Elke Rundensteiner.
KDD, 2019.

6. *Patient-Level Classification of Clinical Note Sequences Guided by Attributed Hierarchical Attention*.
Cansu Sen, **Thomas Hartvigsen**, Xiangnan Kong, Elke Rundensteiner.

IEEE BigData, 2019.

5. *Learning Temporal Relevance in Longitudinal Medical Notes*.
Cansu Sen, **Thomas Hartvigsen**, Xiangnan Kong, Elke Rundensteiner.
IEEE BigData, 2019.

4. *Comparing General and Locally-Learned Word Embeddings for Clinical Text Mining*.
Jidapa Thadajarassiri, Cansu Sen, **Thomas Hartvigsen**, Xiangnan Kong, Elke Rundensteiner.
IEEE BHI, 2019.

3. *Detecting MRSA Infections by Fusing Structured and Unstructured Electronic Health Record Data*.
Thomas Hartvigsen, Cansu Sen, Elke Rundensteiner.
BIOSTEC, 2018.

2. *Early Prediction of MRSA Infections using Electronic Health Records*.
Thomas Hartvigsen, Cansu Sen, Sarah Brownell, Erin Teeple, Xiangnan Kong, Elke Rundensteiner.
HEALTHINF, 2018. 🏆 **Best Student Paper runner up**.

1. *CREST - Risk Prediction for Clostridium Difficile Infection Using Multimodal Data Mining*.
Cansu Sen, **Thomas Hartvigsen**, Kajal Claypool, Elke Rundensteiner.
ECML, 2017.

In-submission

9. TOXIGEN: *Controlling Language Models to Generate Implied and Adversarial Toxicity*.
Thomas Hartvigsen, Saadia Gabriel, Hamid Palangi, Maarten Sap, Dipankar Ray, Ece Kamar.

8. *Continuous-Time Attention Network for Irregularly-Sampled Time Series Classification*.
Thomas Hartvigsen, Jidapa Thadajarassiri, Xiangnan Kong, Elke Rundensteiner.

7. *Stop&Hop: Early Classification of Irregular Time Series*.
Thomas Hartvigsen, Walter Gerych, Jidapa Thadajarassiri, Xiangnan Kong, Elke Rundensteiner.

6. *Class-Specific Explainability for Deep Time Series Classifiers*.
Ramesh Doddaiah, Prathyush Parvatharaju, Elke Rundensteiner, **Thomas Hartvigsen**.

5. *Knowledge Amalgamation for Multi-Label Classification via Label Dependency Transfer*.
Jidapa Thadajarassiri, **Thomas Hartvigsen**, Walter Gerych, Xiangnan Kong, Elke Rundensteiner.

4. *The Road to Explainability is Paved with Bias: Measuring the Fairness of Explanations*.
Aparna Balagopalan, Haoran Zhang, Kimia Hamidieh, **Thomas Hartvigsen**, Frank Rudzicz, Marzyeh Ghassemi.

3. *SAIL: Recurrent Classifier Chains with Incomplete Labels*.
Walter Gerych, **Thomas Hartvigsen**, Emmanuel Agu, Elke Rundensteiner.

2. *SkipSNN: Efficiently Classifying Noisy Spike Trains*.
Hang Yin, Xiangnan Kong, Liping Liu, **Thomas Hartvigsen**, Xin Dai.

1. *Multi-State Brain Network Discovery*.
Hang Yin, Xinyue Liu, Xiangnan Kong, **Thomas Hartvigsen**, Yanhua Li.

Supervised Undergraduate Papers

3. *Early Diagnosis Prediction with Recurrent Neural Networks*.
Daniel Johnston†, Liubou Klindziuk†, Lolita Nazarov†, **Thomas Hartvigsen**, Elke Rundensteiner.
IEEE URTC 2019. 🏆 Best Paper runner up.

2. *Handling Missing Values in Multivariate Time Series Classification*.
Julia Friend†, Alec Hauck†, Sruthi Kurada†, Cansu Sen, **Thomas Hartvigsen**, Elke Rundensteiner.
IEEE URTC 2018.

1. *MRSA Infection Prediction System*.
Sarah Brownell†, **Thomas Hartvigsen**, Elke Rundensteiner.

IEEE URTC 2017.

†undergraduate co-author

TALKS

Timely and Trustworthy Machine Learning
MIT, HealthyML group 11/2021

Adaptive-Halting Policy Networks for Early Classification
Computational Sustainability Doctoral Consortium 10/2020
Harvard University, DtAK group 09/2020
Florida State University Data Science Seminar 06/2020
MITRE Data Science 03/2020
IBM Research, SystemML group 01/2020
University of Minnesota, IMA 09/2019
Northeastern University, NEML Day poster 09/2019
ACM SIGKDD 08/2019

Recurrent Halting Chains for Early Multi-label Classification
ACM SIGKDD 08/2020

Recurrent Models for Clinical Time Series
Worcester Polytechnic Institute, Arts & Sciences Showcase 05/2019

TEACHING/MENTORING

I have supervised one PhD qualifier, two masters theses, and eight undergraduate students. I generally present project ideas to students, then meet with them once or twice a week to resolve any issues.

Students Advised:

- Prathyush Parvatharaju, MS, WPI 2019-now
  - **Masters Thesis**: *Learning Saliency Maps to Explain Deep Time Series Classifiers*
- Ramesh Doddaiah, PhD, WPI 2020-now
  - **PhD Qualifier**: *Class-Specific Explainability for Deep Time Series Classifiers*
- Aleksa Perucic, MS, WPI 2019-2020
  - **Masters Thesis**: *SIFT - A Deep Network for Irregular Multivariate Time Series*
- Liubuo (Yuuna) Klindziuk, BS, Amherst College 2019
- Daniel Johnston, BS, Columbia University 2019
- Lolita Nazarov, BS, StonyBrook University 2019
- Julia Friend, BS, Oberlin College 2018
- Alex Hauck, BS, Valporaiso University 2018
- Sruthi Kurada, Advanced Math & Science Academy Charter School 2018
- Sarah Brownell, BS, Simmons University 2017
- Sean Tocci, BS, UMass Dartmouth 2017

Developed workshop on Deep Learning with PyTorch for Undergrads, WPI. 2019

SERVICE

Conference Program Committee:

- ACL Rolling Reviewer (2021-present)
- AAAI ('21, '22)
- CVPR ('21)
- ICCV ('21)
- ACL ('21, '22)
- EMNLP ('21)

External Reviewer: KDD ('18, '19, '20)
Conference Volunteer: KDD ('19, '20, '21), NeurIPS ('20, '21)
Deep Learning Reading Group, Organizer, WPI 2019-2020
Graduate Student Advisory Council to the Dean of Arts & Sciences, WPI 2018-2020
Graduate Student Government Senate, WPI 2018
Data Science Graduate Student Council, WPI 2016-2019

REFERENCES

Available upon request